

Project 4

Constraint Satisfaction Problems (CSPs)

Deadline: 2020-7-17 Friday 23:59

July 6, 2020

Introduction

Abstract

A constraint satisfaction problem (CSP) is a problem whose solution is an assignment of values to variables that respects the constraints on the assignment and the variables' domains.

CSPs are very powerful because a single fixed set of algorithms can be used to solve any problem specified as a CSP, and many problems can be intuitively specified as CSPs.

The goal of this project is to help you better understand CSPs and how they are solved. In this project you will be implementing a CSP solver and improving it with heuristic and inference algorithms.

Starter Files

For this project, you are provided with an autograder as well as many test cases that specify CSP problems. These test cases are specified within the `csps` directory. The tree diagram of the starter system is as follows.

```
.
+-- BinaryCSP.py
+-- Interface.py
+-- StudentAutograder.py
+-- Testing.py
+-- csps
|   +-- csp7.assignment
|   +-- csp7.csp
|   +-- ...
+-- test_cases
    +-- backtracking3solvable1_1.solution
```

```
+++ backtracking3solvable1_1.test
+++ ...
```

Functions are provided in `Testing` to parse these files into the objects your algorithms will work with. The files follow a simple format you can understand by inspection, so if your code does not work, looking at the problems themselves and manually checking your logic is the best debugging strategy.

Data Structure

All of the necessary interface structures for assignments and CSP problems is provided for you in `Interface.py`, and every function you will be implementing are marked with `TODO` in `BinaryCSP.py`. While you do not need to implement these structures, it is important to understand how they work.

Almost every function you will be implementing will take in

- a `ConstraintSatisfactionProblem`
- an `Assignment`

The `ConstraintSatisfactionProblem` object serves only as a representation of the problem, and it is not intended to be changed. It holds three things:

- a dictionary from variables to their domains (`varDomains`)
- a list of binary constraints (`binaryConstraints`)
- a list of unary constraints (`unaryConstraints`).

An `Assignment` is constructed from a `ConstraintSatisfactionProblem` and is intended to be updated as you search for the solution. It holds

- a dictionary from variables to their domains (`varDomains`)
- a dictionary from variables to their assigned values (`assignedValues`).

Notice that the `varDomains` in `Assignment` is meant to be updated, while the `varDomains` in `ConstraintSatisfactionProblem` should be left alone.

A new assignment should never be created. All changes to the assignment through the recursive backtracking and the inference methods that you will be implementing are designed to be reversible. This prevents the need to create multiple assignment objects, which becomes very space-consuming.

The constraints in the CSP are represented by two classes:

- `BinaryConstraint`
- `UnaryConstraint`

Both of these store the variables affected and have an `isSatisfied` function that takes in the value(s) and returns `False` if the constraint is broken. You will only be working with binary constraints, as `eliminateUnaryConstraints` has been implemented of for you.

Two useful methods for binary constraints include `affects`, which takes in a variable and returns `True` if the constraint has any impact on the variable, and `otherVariable`, which takes in one variable of the `BinaryConstraint` and returns the other variable affected.

Questions

Question 1: Recursive Backtracking

In this question you will be creating the basic **recursive backtracking framework** for solving a constraint satisfaction problem.

First, implement the function `consistent`.

```
def consistent(assignment, csp, var, value):
    """
    Checks if a value assigned to a variable is consistent with all binary
    constraints in a problem.

    * Do not assign value to var.
    * Only check if this value would be consistent or not.
    * If the other variable for a constraint is not assigned, then the new
      value is consistent with the constraint.

    Args:
    * assignment (Assignment): the partial assignment
    * csp (ConstraintSatisfactionProblem): the problem definition
    * var (string): the variable that would be assigned
    * value (value): the value that would be assigned to the variable

    Returns:
    * boolean
      True if the value would be consistent with all currently assigned values,
      False otherwise.
    """
```

- This function indicates whether a given value would be possible to assign to a variable without violating any of its constraints.
- You only need to consider the constraints in `csp.binaryConstraints` that affect this variable and have the other affected variable already assigned.

Once this is done, implement `recursiveBacktracking`

```
def recursiveBacktracking(assignment, csp, orderValuesMethod,
                          selectVariableMethod, inferenceMethod):
    """
    Recursive backtracking algorithm.

    * A new assignment should not be created.
    * The assignment passed in should have its domains updated with inferences.
    * In the case that a recursive call returns failure or a variable assignment
      is incorrect, the inferences made along the way should be reversed.
    * See maintainArcConsistency and forwardChecking for the format of inferences.

    Examples of the functions to be passed in:
    * orderValuesMethod: orderValues, leastConstrainingValuesHeuristic
    * selectVariableMethod: chooseFirstVariable, minimumRemainingValuesHeuristic
    * inferenceMethod: noInferences, maintainArcConsistency, forwardChecking

    Args:
    * assignment (Assignment): a partial assignment to expand upon
    * csp (ConstraintSatisfactionProblem): the problem definition
    * orderValuesMethod (function<assignment, csp, variable> returns list<value>):
      a function to decide the next value to try
    * selectVariableMethod (function<assignment, csp> returns variable):
      a function to decide which variable to assign next
    * inferenceMethod (function<assignment, csp, variable, value> returns
      set<variable, value>):
      a function to specify what type of inferences to use

    Returns:
    * Assignment
      A completed and consistent assignment. None if no solution exists.
    """
```

The algorithm is provided as follows, see also Slide 26 of Lecture 12, CSP I.

```
function BACKTRACKING-SEARCH(csp) returns solution/failure
  return RECURSIVE-BACKTRACKING({ }, csp)

function RECURSIVE-BACKTRACKING(assignment, csp) returns soln/failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(VARIABLES[csp], assignment, csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment given CONSTRAINTS[csp] then
      add {var = value} to assignment
      result ← RECURSIVE-BACKTRACKING(assignment, csp)
      if result ≠ failure then return result
      remove {var = value} from assignment
  return failure
```

- Ignore for now the parameter `inferenceMethod`.
- This function is designed to take in a problem definition and a partial assignment. When finished, the assignment should either be a complete solution to the CSP or indicate failure.

Question 2: Variable Selection

While the recursive backtracking method eventually finds a solution for a constraint satisfaction problem, this basic solution will take a very long time for larger problems.

Fortunately, there are heuristics that can be used to make it faster, and one place to include heuristics is in selecting which variable to consider next.

Implement `minimumRemainingValueHeuristic`.

```
def minimumRemainingValuesHeuristic(assignment, csp):
    """
    Selects the next variable to try to give a value to in an assignment.
    * Uses minimum remaining values heuristic to pick a variable. Use degree
    heuristic for breaking ties.

    Args:
    * assignment (Assignment): the partial assignment to expand
    * csp (ConstraintSatisfactionProblem): the problem description

    Returns:
    * the next variable to assign
    """
```

- This follows the minimum remaining value heuristic (see Slide 41 of CSP I) to select a variable with the fewest options left and uses the degree heuristic to break ties.
- The degree heuristic chooses the variable that is involved in the largest number of constraints on other unassigned variables.

Question 3: Value Ordering

Another way to use heuristics is to optimize a constraint satisfaction problem solver is to attempt values in a different order.

Implement `leastConstrainingValuesHeuristic`.

```
def leastConstrainingValuesHeuristic(assignment, csp, var):
    """
    Creates an ordered list of the remaining values left for a given variable.
    * Values should be attempted in the order returned.
    * The least constraining value should be at the front of the list.

    Args:
    * assignment (Assignment): the partial assignment to expand
    * csp (ConstraintSatisfactionProblem): the problem description
    * var (string): the variable to be assigned the values

    Returns:
    * list<values>
        a list of the possible values ordered by the least constraining value
        heuristic
    """
```

- This takes in a variable and determines the order in which the possible values should be attempted according to the least constraining values heuristic (see Slide 42 of CSP I), which prefers values that eliminate the fewest possibilities from other variables.
- Your code should be able to solve small CSPs. To test and debug, use `StudentAutograder.py` and the functions in `Testing.py`.

Question 4 (4 points): Forward Checking

While heuristics help to determine what to attempt next, there are times when a particular search path is doomed to fail long before all of the values have been tried.

Inferences (see Slides 28-36 of CSP I) are a way to identify impossible assignments early on by looking at how a new value assignment affects other variables.

The pseudocode of `recursiveBacktracking` that integrates inferences is given.

```
function BACKTRACKING-SEARCH(csp) returns a solution, or failure
  return BACKTRACK({ }, csp)

function BACKTRACK(assignment, csp) returns a solution, or failure
  if assignment is complete then return assignment
  var ← SELECT-UNASSIGNED-VARIABLE(csp)
  for each value in ORDER-DOMAIN-VALUES(var, assignment, csp) do
    if value is consistent with assignment then
      add {var = value} to assignment
      inferences ← INFERENCE(csp, var, value)
      if inferences ≠ failure then
        add inferences to assignment
        result ← BACKTRACK(assignment, csp)
        if result ≠ failure then
          return result
      remove {var = value} and inferences from assignment
  return failure
```

Figure 6.5 A simple backtracking algorithm for constraint satisfaction problems. The algorithm is modeled on the recursive depth-first search of Chapter 3. By varying the functions `SELECT-UNASSIGNED-VARIABLE` and `ORDER-DOMAIN-VALUES`, we can implement the general-purpose heuristics discussed in the text. The function `INFERENCE` can optionally be used to impose arc-, path-, or k -consistency, as desired. If a value choice leads to failure (noticed either by `INFERENCE` or by `BACKTRACK`), then value assignments (including those made by `INFERENCE`) are removed from the current assignment and a new value is tried.

Each inference made by an algorithm involves one possible value being removed from one variable. It should be noted that when these inferences are made they must be kept track of so that they can later be reversed if a particular assignment fails. They are stored as a set of tuples (variable, value).

First, update `recursiveBacktracking` to deal with the parameter `inferenceMethod`.

Then, implement `forwardChecking`.

```
def forwardChecking(assignment, csp, var, value):
    """
    Implements the forward checking algorithm.

    * Each inference should take the form of (variable, value) where the value
    is being removed from the domain of variable.
```

- * This format is important so that the inferences can be reversed if they result in a conflicting partial assignment.
- * If the algorithm reveals an inconsistency, any inferences made should be reversed before ending the function.

Args:

- * assignment (Assignment): the partial assignment to expand
- * csp (ConstraintSatisfactionProblem): the problem description
- * var (string): the variable that has just been assigned a value
- * value (string): the value that has just been assigned

Returns:

- * set< tuple<variable, value> >
the inferences made in this call or None if inconsistent assignment

"""

This is a very basic `inferenceMethod` function.

- When a value is assigned, all variables connected to the assigned variable by a binary constraint are considered.
- If any value in those variables is inconsistent with that constraint and the newly assigned value, then the inconsistent value is removed.
- You can test again your code by passing `forwardChecking` to `inferenceMethod`. **It should be faster than the previous version in Question 3.**

Question 5: Maintaining Arc Consistency

There are other methods for making inferences than can detect inconsistencies earlier than forward checking. One of these is the Maintaining Arc Consistency algorithm, or the MAC algorithm.

It works like the AC-3 algorithm (see Slide 35 of CSP I). We recall a version of AC-3 where inferences are saved.

```

function AC-3(csp) returns false if an inconsistency is found and true otherwise
inputs: csp, a binary CSP with components ( $X$ ,  $D$ ,  $C$ )
local variables: queue, a queue of arcs, initially all the arcs in csp

while queue is not empty do
    ( $X_i$ ,  $X_j$ )  $\leftarrow$  REMOVE-FIRST(queue)
    if REVISE(csp,  $X_i$ ,  $X_j$ ) then
        if size of  $D_i$  = 0 then return false
        for each  $X_k$  in  $X_i$ .NEIGHBORS -  $\{X_j\}$  do
            add ( $X_k$ ,  $X_i$ ) to queue
return true

```



```

function REVISE(csp,  $X_i$ ,  $X_j$ ) returns true iff we revise the domain of  $X_i$ 
    revised  $\leftarrow$  false
    for each  $x$  in  $D_i$  do
        if no value  $y$  in  $D_j$  allows  $(x,y)$  to satisfy the constraint between  $X_i$  and  $X_j$  then
            delete  $x$  from  $D_i$ 
            revised  $\leftarrow$  true
    return revised

```

Figure 6.3 The arc-consistency algorithm AC-3. After applying AC-3, either every arc is arc-consistent, or some variable has an empty domain, indicating that the CSP cannot be solved. The name “AC-3” was used by the algorithm’s inventor (Mackworth, 1977) because it’s the third version developed in the paper.

The difference between MAC and AC-3 are as follows.

- While AC-3 is called as a preprocessing step, which explains why all arcs inserted in the queue, MAC is called during the search.
- After a variable X_i is assigned a value, MAC performs the same operations as AC-3, but starts with only the arcs (X_j, X_i) for all X_j that are unassigned variables that are neighbors of X_i .

First, implement revise.

```

def revise(assignment, csp, var1, var2, constraint):
    """
    Helper function to maintainArcConsistency and AC3.

    * Remove values from var2 domain if constraint cannot be satisfied.
    * Each inference should take the form of (variable, value) where the value
      is being removed from the domain of variable.
    * This format is important so that the inferences can be reversed if they
      result in a conflicting partial assignment.
    * If the algorithm reveals an inconsistency, any inferences made should be
      reversed before ending the function.

    Args:
    * assignment (Assignment): the partial assignment to expand
    * csp (ConstraintSatisfactionProblem): the problem description
    * var1 (string): the variable with consistent values
    * var2 (string): the variable that should have inconsistent values removed
    * constraint (BinaryConstraint): the constraint connecting var1 and var2

    Returns:
    * set<tuple<variable, value>>
      the inferences made in this call or None if inconsistent assignment
    """

```

This is a helper function that is responsible for determining inconsistent values in a variable. Then implement `maintainArcConsistency`.

```
def maintainArcConsistency(assignment, csp, var, value):
    """
    Implements the maintaining arc consistency algorithm.

    * Inferences take the form of (variable, value) where the value
    is being removed from the domain of variable.
    * This format is important so that the inferences can be reversed if they
    result in a conflicting partial assignment.
    * If the algorithm reveals an inconsistency, and inferences made should be
    reversed before ending the function.

    Args:
    * assignment (Assignment): the partial assignment to expand
    * csp (ConstraintSatisfactionProblem): the problem description
    * var (string): the variable that has just been assigned a value
    * value (string): the value that has just been assigned

    Returns:
    * set<<variable, value>>
      the inferences made in this call or None if inconsistent assignment
    """
```

The MAC algorithm starts off very similar to forward checking in that it removes inconsistent values from variables connected to the newly assigned variable.

The difference is that it uses a **queue** to propagate these changes to other related variables.

Question 6: Preprocessing

Another step to making a constraint satisfaction solver more efficient is to perform preprocessing. This can eliminate impossible values before the recursive backtracking even starts.

One method to do this is to use the AC-3 algorithm.

Implement AC3.

```
def AC3(assignment, csp):
    """
    AC3 algorithm for constraint propagation.

    * Used as a pre-processing step to reduce the problem before running
    recursive backtracking.
```

```

Args:
* assignment (Assignment): the partial assignment to expand
* csp (ConstraintSatisfactionProblem): the problem description

Returns:
* Assignment
    the updated assignment after inferences are made or None if an
    inconsistent assignment
"""

```

Note it does not need to track the inferences that are made, because if the assignment fails at any point then there is no prior state to back up to. This means that there is no solution to the CSP.

Grading

Submit `BinaryCSP.py` in a zip file to online judge.

Submission The `BinaryCSP.py` is where your entire CSP implementation will reside. You should submit this file (and only this one) with your code and comments to online judge.

Evaluation Your code will be autograded for technical correctness. However, the correctness of your implementation, not the autograder's judgements, will be the final judge of your score. If necessary, we will review and grade assignments individually to ensure that you receive due credit for your work.

Testcases The information of testcases are given as follows.

Question ID	Score per case	Public Case ID	Hidden Case ID
1	4	1 - 7	8 - 10
2	2	11 - 17	18 - 20
3	2	21 - 24	25 - 27
4	4	28 - 35	36 - 38
5	4	39 - 45	46 - 49
6	2	50 - 57	57 - 62

Some testcases have been provided in the starter for you to test. Please use `StudentAutograder.py` to test your program, and do not use the online judge as a method of debugging.

Honor Code We will be checking your code against other submissions in the class for logical redundancy. If you copy someone else's code from online and submit it with minor changes, we will know. These cheat detectors are quite hard to fool, so please don't try.